

# L5: HPC for AI applications & Environmental impact of computation

---

P. de Oliveira Castro, M. Jam

October 2, 2026

Master Calcul Haute Performance et Simulation - GLHPC | UVSQ

- 1 HPC for AI & Environmental impact of computation
- 2 Introduction to AI applications
- 3 Environmental impact of computation
- 4 Energy consumption of HPC
- 5 AI energy and computation costs
- 6 More frugal computing?

## HPC for AI & Environmental impact of computation

---

## Introduction to AI applications

---

- 2012: **AI renaissance** brought by increased data availability and computation resources
  - breakthroughs in multiple domains
  - many innovations: algorithms, specialized processors, optimizations
- Most systems use **neural networks**:
  - Training (stochastic gradient descent + backpropagation)
  - Inference (forward pass)
- For both, the **bottleneck** is matrix multiplication

# Objectives

- Explain why dense linear algebra (GEMM) dominates NN compute
- Core SGEMM kernel ideas and common optimizations
- Use Roofline model to identify bottlenecks

## Short introduction to Neural Networks

- Neural networks are composed of layers of neurons
- Each neuron computes a weighted sum of its inputs followed by a non-linear activation function  $f$

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \rightarrow \text{neuron} \rightarrow y$$


$$y = f \left( \sum_i w_i x_i + b \right)$$

- Common activation functions: ReLU, sigmoid, ...
- Perceptron: single layer of neurons (1958 Rosenblatt)

# Architectures

- Different architectures for different tasks:
  - Fully connected layers
  - Convolutional layers
  - Recursive layers
  - Transformers (attention mechanism)

Inputlayer



- Inference: use the trained model to make predictions on new data
- Forward pass through the network:
  - For each layer, compute the weighted sum and apply activation function
  - The weighted sum is a matrix-vector multiplication for fully connected layers and convolutions (often implemented as GEMM).

## Two layer network

Layer 1:

- $X$ : input data  $[K \times B] \rightarrow K$  features,  $B$  batch size
- $W_1$ : weights  $[H \times K] \rightarrow H$  hidden units
- $b_1$ : bias  $[H \times 1]$  (broadcasted to batch  $B$ )

Layer 2:

- $W_2$ : weights  $[O \times H] \rightarrow O$  outputs
- $b_2$ : bias  $[O \times 1]$  (broadcasted to batch  $B$ )

ReLU  $f(x) = \max(0, x)$ ,  $f'(x) = 1_{x>0}$

## Forward inference

- Layer 1 Pre-activation hidden (GEMM,  $H \times K \times K \times B \rightarrow H \times B$ )

$$Z_1 = W_1 \cdot X + b_1$$

- Layer 1 Activation - ReLU (elementwise)

$$H = f(Z_1)$$

- Layer 2 Output pre-activation (GEMM,  $O \times H \times H \times B \rightarrow O \times B$ )

$$Z_2 = W_2 \cdot H + b_2$$

- Layer 2 Activation - ReLU (elementwise)

$$Y = f(Z_2)$$

- Forward is dominated by the two large GEMMs  $Z_1$  and  $Z_2$ .

# Training

- Training: adjust weights  $W$  and biases  $b$  to minimize a loss function  $L$  over a training dataset
- Use **backpropagation** to compute gradients on each layer (chain rule)
- Example with one neuron and MSE loss:

$$z = w_1x_1 + w_2x_2 + b, \quad y = f(z)$$

$$L = (y - y_{true})^2$$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial y} \cdot \frac{\partial y}{\partial z} \cdot \frac{\partial z}{\partial w_1} = 2(y - y_{true}) \cdot f'(z) \cdot x_1$$

- Backward pass can be efficiently implemented using automatic differentiation and matrix multiplications.

- Use stochastic gradient descent to update weights:

$$w_1 \leftarrow w_1 - \eta \cdot \frac{\partial L}{\partial w_1}$$

$$w_2 \leftarrow w_2 - \eta \cdot \frac{\partial L}{\partial w_2}$$

$$b \leftarrow b - \eta \cdot \frac{\partial L}{\partial b}$$

- $\eta$  is the learning rate
- Repeat for many epochs over the training dataset

1. Forward pass to compute  $H$  and  $Y$
2. Compute loss  $L(Y, Y_{true})$
3. Backward pass to compute gradients.

The backward pass is also dominated by GEMMs.

- Popular frameworks: TensorFlow, PyTorch, JAX, ...
- High-level APIs for defining models, automatic differentiation, GPU acceleration

```
# Simple 2-layer NN in PyTorch
import torch
import torch.nn as nn


class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(28*28, 512)
        self.fc2 = nn.Linear(512, 10)

    def forward(self, x):
        x = torch.flatten(x, 1)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        return x
```

# SGEMM

Single-precision General Matrix-Matrix multiplication (SGEMM):

$$RES = A \times B + C$$



## Naive SGEMM implementation (pseudocode)

```
// Initialize RES to C
for (i = 0; i < M; i++)
    for (j = 0; j < N; j++)
        RES[i][j] = C[i][j];

// Matrix multiply
for (i = 0; i < M; i++) {
    for (j = 0; j < N; j++) {
        for (k = 0; k < K; k++) {
            RES[i][j] += A[i][k] * B[k][j];
        }
    }
}
```

- FLOPS:  $2 \times M \times N \times K$
- min. Memory: 4 bytes  $\times (M \times K + K \times N + M \times N)$

order in memory  $\rightarrow$

$$\begin{bmatrix} b_{11} & b_{12} & b_{13} & b_{14} \\ b_{21} & b_{22} & b_{23} & b_{24} \\ b_{31} & b_{32} & b_{33} & b_{34} \\ b_{41} & b_{42} & b_{43} & b_{44} \end{bmatrix}$$

- Stride in accessing B (column-major)
  - Poor spatial locality
  - Difficult to vectorize
  - Cache misses for large matrices (reuse distance too large)
- **Low arithmetic intensity:**  $\approx 0.5$  FLOP/byte for large matrices

## Reordering loops (i,k,j)

- Sums  $RES[i][j] += A[i][k] * B[k][j]$ ; are independent  $\rightarrow$  reorder loops:

```
for (i = 0; i < M; i++)  
  for (k = 0; k < K; k++)  
    for (j = 0; j < N; j++)  
      RES[i][j] += A[i][k] * B[k][j];
```

- $A[i][k]$  does not depend on  $j \rightarrow$  load once, reuse  $N$  times
- $RES$  and  $B$  accesses are now stride-1 (row-major)

```
for (i = 0; i < M; i++)  
  for (k = 0; k < K; k++) {  
    const float temp = A[i][k];  
    for (j = 0; j < N; j++)  
      RES[i][j] += temp * B[k][j];  
  }
```

- Better spatial locality and easier to vectorize

Inner loop assembly for (i,k,j) ordering with AVX (8 float in a vector):

```
.loop:                                # Inner loop
    vmovss  xmm0, DWORD PTR A[i][k]  # Load A[i][k]
    vbroadcastss ymm0, xmm0          # Broadcast scalar to
    all lanes
    vmovaps ymm1, YMMWORD PTR B[k][j] # Load B[k][j:j+8]
    vfmadd231ps ymm2, ymm1, ymm0      # Fused multiply-add
    vmovaps YMMWORD PTR RES[i][j], ymm2 # Store RES[i][j:j+8]
    add     j, 8                      # Increment j by 8 (
    vector width)
    cmp     j, N                      # Compare j with N
    jl     .loop                     # Loop if j < N
```

- Temporal locality analysis:
  - GOOD:  $A[i][k]$  reused in the inner loop, reuse distance 1.
  - MEDIUM: For a given  $(i, j)$ , each  $RES[i][j]$  revisited once per  $k$ . So reuse distance  $N$  (one full row).
    - To keep RES in cache between uses you would need cache  $\geq N \times 4B$
  - BAD: For a given  $(k, j)$ ,  $B[k][j]$  used once per  $i$ . So reuse distance  $K \times N$  (entire B matrix).
    - To keep B in cache between uses you would need cache  $\geq K \times N \times 4B$
- Still poor temporal locality for large matrices
- Solution: tiling / blocking to increase reuse

## Blocking (tiling)

- **Idea:** operate on sub-matrices blocks that fit in cache

$$\begin{bmatrix} \textcolor{red}{A}_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} \textcolor{blue}{B}_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} \textcolor{red}{A}_{11}\textcolor{blue}{B}_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}\textcolor{blue}{B}_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \end{bmatrix}$$

```
#define BS 64 // Block size
// Loop over blocks
for (ii = 0; ii < M; ii += BS)
    for (kk = 0; kk < K; kk += BS)
        for (jj = 0; jj < N; jj += BS)

            // Operate on blocks A[ii:ii+BS, kk:kk+BS],
            // B[kk:kk+BS, jj:jj+BS], RES[ii:ii+BS, jj:jj+BS]
            for (i = ii; i < min(ii+BS, M); i++)
                for (k = kk; k < min(kk+BS, K); k++)
                    for (j = jj; j < min(jj+BS, N); j++)
                        RES[i][j] += A[i][k] * B[k][j];
```

- Each block operation is independent → parallelize over blocks


```
#pragma omp parallel for collapse(3)
for (ii = 0; ii < M; ii += BS)
    for (jj = 0; jj < N; jj += BS)
        for (kk = 0; kk < K; kk += BS)
            // Block multiplication as before
```

- Each thread works on its own block → no false sharing
- Synchronization only at the end of the parallel region
- NUMA considerations: pin threads to cores, allocate memory close to threads
- Load balancing: static scheduling usually works well for large matrices


- Highly optimized SGEMM implementations exist:
  - OpenBLAS, Intel MKL for CPU
  - NVIDIA cuBLAS for GPU
- Implementations use blocking, vectorization, parallelization, and many architecture-specific optimizations
- Libraries are carefully tuned for different sizes and shapes of matrices.
- Autotuners (e.g., ATLAS, TVM, MLKAPS) can generate optimized code for specific hardware and problem sizes.

- Hypothesis: performance is limited by either compute or memory bandwidth
  - performance: FLOP/s (vertical axis)
  - memory bandwidth: Bytes/s
  - arithmetic intensity: FLOP/byte (horizontal axis)
- Simple visual model to understand bottlenecks

## Roofline model - Bounds



# Roofline model - SGEMM analysis



## Environmental impact of computation

---

- **Major ecological crisis:** French roadmap targets carbon neutrality in 2050 (Stratégie Nationale Bas Carbone).
- Requires a **40% energy consumption reduction**.
- **HPC part of the solution:** modeling and improving complex systems

## HPC part of the problem

- Frontier system at ORNL
  - More than  $10^{18}$  floating point operations per second
  - Consumes 21MW: the energy of a small town (16 000 french houses)



- The ICT sector consumes  $\approx 5\%$  of the energy worldwide
- It accounts for **1.8% - 2.8%** of emitted GHG [Freitag, 2021]:
  - Accounts for embodied emissions.
  - Shadow energy during the **whole life-cycle**: mining, fabrication, transportation, recycling.
- GHG emissions are only one of the sustainability issues
  - rare-earth mining and waste disposal (eg. Agbogbloshie).
    - human-right abuses, health issues, pollution.
- This presentation focus on energy consumption of HPC


## What about renewable energies?

- Low-carbon electricity is a **limited** resource
- Decarbonation → huge increase in electricity demand
  - Heating, Transportation, Industry
  - Computing will compete for low-carbon electricity.

## Energy consumption of HPC

---

# Evolution of processing units [Batten, 2023]



$$\text{CMOS Power } P = \underbrace{1/2 \cdot C \cdot V^2 \cdot f}_{P_{\text{dynamic}}} + \underbrace{V \cdot I_{\text{leak}}}_{P_{\text{static}}}$$

For each generation, transistors dimensions reduced by 30%,

- Voltage and capacitance reduced by 30%
- Frequency increases:  $\times 1.4 \approx 1/0.7$
- Surface halved:  $0.5 \approx 0.7 \times 0.7$
- Power halved:  $\Delta P = 0.7 \times 0.7^2 \times 1/0.7 \approx 0.5$

**Power per surface unit remains constant** but manufacturers double number of transistors and frequency increases:

- Power efficiency doubles every 1.57 years
- Total power increases

- At current scale, leak currents start increasing ( $P_{\text{static}} \nearrow$ ).  
Power wall slows Dennard's scaling.
- Computing demand  $\rightarrow$  parallelism and specialization.
- Number of cores increases exponentially since 2005.
- Power efficiency still improving:
  - selectively turning-off inactive transistors;
  - architecture design optimizations;
  - software optimizations.

- For domain specific applications, such as AI, specialized accelerators are used
  - Memory and compute units tuned for a specific problem (matrix multiplication) ;
  - Faster and better power efficiency: GPU, TPU, FPGA, ASIC.

# Analysis of TOP-100 HPC systems




Figure 7: Evolution of TOP 100 systems

Efficiency and Peak computation exponential increase.

- In 1865, Jevons shows that steam engine improvements translate into increased coal consumption.
- In HPC, efficiency gains contribute to the rising computation demand.
  - net increase of the total power consumption.
- Rebound effects for data-centers [Masanet, 2020]
  - 6% increase in energy consumption from 2010 to 2018 (255 % increase in nodes).
- **Indirect rebound effects:** computation advances can contribute to the acceleration of other fields.

## AI energy and computation costs

---

# Training cost doubles every 3.4 months [OpenAI, 2020]





Figure 8: OpenAI, 2020

# Should we study training or inference?

- Early industry rule-of-thumb: **80–90%** of lifecycle cost is inference [NVIDIA, 2019]
  - Assumed a single model trained once and served to billions of queries
- Real hyperscaler measurements show **training remains a major fraction**:
  - Google [Patterson et al., 2022]: **~40% training vs. ~60% inference** (3:2 ratio) across all ML workloads
  - Meta [Wu et al., MLSys 2022]: **29% training, 40% inference, 31% data processing**
- **Frontier models keep training compute massive**:
  - Training is not “done once”: thousands of experimentation runs, ablations, synthetic data generation, RLHF
  - Llama 3.1 405B:  $\approx 3.8 \times 10^{25}$  FLOPs ( $\approx 39$  GWh on 16k H100 GPUs) [Meta, 2024]
- **Both matter**: training drives extreme-scale HPC clusters (10k–100k GPUs, interconnects), while inference dominates continuous 24/7 energy

# Inference cost - Diminishing returns for computer vision



Exponential increase in compute for linear accuracy gain  
[Desislavov, 2023 / Schwartz, 2019]

More frugal computing?

---

## Smaller precision / Smaller models for AI





Figure 9: Shankar 2022

- Inference cost grows with model complexity
- Simpler models are often more interpretable
  - Traditional science also prefers simpler models
- DNN not necessary for all tasks

- Knights Mill 72 cores
- Intel MKL dgetrf
- $n \in [1000, 3000]$
- RAPL estimation


(Thomas Roglin, M1 UVSQ/INTEL internship 2023)

## Save energy by computing slower: 1GHz



- Model: RAPL + 40W
- System power dominates at low frequencies

## Race to idle: 2.6 GHz compute faster and turn off machine



# Need for an interdisciplinary discussion

- AI / HPC can contribute towards sustainability (eg. acceleration of weather forecast models) ... **but its energy cost must be reduced**
- Efficiency:
  - Improve hardware and software
  - Use smaller models / smaller precision
- ... **subject to rebound effects**
- Frugality in computing:
  - Balance computation cost vs. outcomes for each task
  - Choose the right sized model
  - Assess the environmental impact

Treatment of febrile children illnesses in dispensaries.

- IMCI: Paper-based decision tree WHO
- e-POCT CART tree tailored to real data on a standalone tablet
  - Final CART tree easy to interpret and manually checked
  - Randomized-trial → better clinical outcomes and antibiotic prescription reduction
- Sophisticated AI that continuously collects patient data and adapts the algorithm ?
  - Increase in hardware and computation costs.
  - Loss in explainability and verification of the algorithm.

- S. Boehm Optimizing, How to Optimize a CUDA Matmul Kernel

## References - Environmental impact of computation

- Jones, Nicola (2018) 'How to stop data centres from gobbling up the world's electricity'. *Nature*, 561(7722), pp. 163–167.
- Freitag, Charlotte, Berners-Lee, Mike, Widdicks, Kelly, Knowles, Bran, et al. (2021) 'The real climate and transformative impact of ICT: A critique of estimates, trends, and regulations'. *Patterns*, 2(9), p. 100340. online
- Masanet, Eric, Shehabi, Arman, Lei, Nuoa, Smith, Sarah and Koomey, Jonathan (2020) 'Recalibrating global data center energy-use estimates'. *Science*, 367(6481), pp. 984–986.
- Schwartz, Roy, Dodge, Jesse, Smith, Noah A. and Etzioni, Oren (2019) 'Green AI'. arXiv:1907.10597
- Amodei, Dario, Hernandez, Danny, Sastry, Girish, Clark, Jack, et al. (2018) 'AI and compute. OpenAI'.  
<https://openai.com/blog/ai-and-compute/>
- Patterson, David, et al. (2022) 'The Carbon Footprint of Machine Learning Training Will Plateau, Then Shrink'. *IEEE Computer*, 55(7), pp. 18–28.